

In [1]:

```
import pandas as pd
```

In [2]:

```
df = pd.read_csv("/Users/bruce/Documents/Study/phishing ML/Supportive_Docs/my_final_
```

In [3]:

```
df.isnull().sum()
```

Out[3]:

```
url          0
label        0
result       0
dtype: int64
```

In [4]:

```
print(df['label'].value_counts())
print(df['result'].value_counts())
```

```
malicious    108523
benign        35378
Name: label, dtype: int64
1      108523
0       35378
Name: result, dtype: int64
```

In [5]:

```
from urllib.parse import urlparse
```

In [6]:

```
df_process = df
```

In [7]:

```

of URL
url_length'] = df_process['url'].apply(lambda x: len(x))

of Hostname
hostname_length'] = df_process['url'].apply(lambda x: len(str(urlparse(x).hostname)))

length of the path from URL
path_length'] = df_process['url'].apply(lambda x: len(urlparse(x).path))

length of first directory
def fd(URL):
    a = urlparse(URL).path
    if len(a.split('/')[2]) == '///':
        return len(a.split('/')[2])
    else:
        return len(a.split('/')[1])
    return 0

firstDir_length'] = df_process['url'].apply(lambda v: fd(v))

def xyz():
    pass

hyphen_count'] = df_process['url'].apply(lambda i: i.count('-'))

at_sign_count'] = df_process['url'].apply(lambda i: i.count('@'))

```

```
total?''] = df_process['url'].apply(lambda i: i.count('?'))
```

```
total%''] = df_process['url'].apply(lambda i: i.count('%'))
```

```
total.''] = df_process['url'].apply(lambda i: i.count('.'))
```

```
total=''] = df_process['url'].apply(lambda i: i.count('='))
```

```
totalhttp'] = df_process['url'].apply(lambda i: i.count('http'))
```

```
totalhttps'] = df_process['url'].apply(lambda i: i.count('https'))
```

```
totalwww'] = df_process['url'].apply(lambda i: i.count('www'))
```

```
# digits in the URL
```

```
def totaldigits(url):
```

```
    total:
```

```
    isnumeric():
```

```
    totaldigits = totaldigits + 1
```

```
    return totaldigits
```

```
totaldigits'] = df_process['url'].apply(lambda i: totaldigits(i))
```

4/24

• ()

```
label'].value_counts()
```

```
pd.read_csv("URL_test.csv")
```

```
[ 'aaa', 'aarp', 'abarth', 'abb', 'abbott', 'abbvie', 'abc', 'able', 'a
bogado', 'abudhabi', 'ac', 'academy', 'accenture', 'accountant', 'acco
untants', 'aco', 'actor', 'ad', 'adac', 'ads', 'adult', 'ae', 'aeg',
'aero', 'aetna', 'af', 'afl', 'africa', 'ag', 'agakhan', 'agency', 'a
i', 'aig', 'airbus', 'airforce', 'airtel', 'akdn', 'al', 'alfaromeo',
'alibaba', 'alipay', 'allfinanz', 'allstate', 'ally', 'alsace', 'alsto
m', 'am', 'amazon', 'americanexpress', 'americanfamily', 'amex', 'amfa
m', 'amica', 'amsterdam', 'analytics', 'android', 'anquan', 'anz', 'a
o', 'aol', 'apartments', 'app', 'apple', 'aq', 'aquarelle', 'ar', 'ara
b', 'aramco', 'archi', 'army', 'arpa', 'art', 'arte', 'as', 'asda', 'a
sia', 'associates', 'at', 'athleta', 'attorney', 'au', 'auction', 'aud
i', 'audible', 'audio', 'auspost', 'author', 'auto', 'autos', 'avianc
a', 'aw', 'aws', 'ax', 'axa', 'az', 'azure', 'ba', 'baby', 'baidu', 'b
anamex', 'bananarepublic', 'band', 'bank', 'bar', 'barcelona', 'barcla
ycard', 'barclays', 'barefoot', 'bargains', 'baseball', 'basketball',
'bauhaus', 'bayern', 'bb', 'bbc', 'bbt', 'bbva', 'bcg', 'bcn', 'bd',
'be', 'beats', 'beauty', 'beer', 'bentley', 'berlin', 'best', 'bestbu
y', 'bet', 'bf', 'bg', 'bh', 'bharti', 'bi', 'bible', 'bid', 'bike',
'bing', 'bingo', 'bio', 'biz', 'bj', 'black', 'blackfriday', 'blockbus
```

In [8]:

```
df_process.head()
```

Out[8]:

	url	label	result	url_length	hostname_length	path_lengt
0	http://1337x.to/torrent/1048648/American-Snipe...	benign	0	83	8	6
1	http://1337x.to/torrent/1110018/Blackhat-2015-...	benign	0	83	8	6
2	http://1337x.to/torrent/1122940/Blackhat-2015-...	benign	0	83	8	6
3	http://1337x.to/torrent/1124395/Fast-and-Furio...	benign	0	83	8	6
4	http://1337x.to/torrent/1145504/Avengers-Age-o...	benign	0	83	8	6

5 rows × 7 columns

In [9]:

```
df_train= pd.read_csv("URL_phase1.csv")
```

In [10]:

```
df_train.head()
```

Out[10]:

	Unnamed: 0	url	label	result	url_length	hostname_length	path_length
0	0	https://www.google.com	benign	0	22	14	0
1	1	https://www.youtube.com	benign	0	23	15	0
2	2	https://www.facebook.com	benign	0	24	16	0
3	3	https://www.baidu.com	benign	0	21	13	0
4	4	https://www.wikipedia.org	benign	0	25	17	0

5 rows × 24 columns

In [11]:

```
# model training
```

In [12]:

```
# finalize train dataset
```

In [13]:

```
df_train.head()
```

Out[13]:

	Unnamed: 0	url	label	result	url_length	hostname_length	path_length
0	0	https://www.google.com	benign	0	22	14	0
1	1	https://www.youtube.com	benign	0	23	15	0
2	2	https://www.facebook.com	benign	0	24	16	0
3	3	https://www.baidu.com	benign	0	21	13	0
4	4	https://www.wikipedia.org	benign	0	25	17	0

5 rows × 24 columns

In [14]:

```
df_ml_train = df_train
```

In [15]:

```
df_ml_train.drop(['Unnamed: 0', 'url', 'label'], axis=1, inplace=True)
```

In [16]:

```
df_ml_train.head()
```

Out[16]:

	result	url_length	hostname_length	path_length	FirstDir_length	total-	total@	total?	total%
0	0	22	14	0	0	0	0	0	0
1	0	23	15	0	0	0	0	0	0
2	0	24	16	0	0	0	0	0	0
3	0	21	13	0	0	0	0	0	0
4	0	25	17	0	0	0	0	0	0

5 rows × 21 columns

In [17]:

```
df_ml_train_X = df_ml_train.loc[:, df_ml_train.columns != 'result']
```

In [18]:

```
df_ml_train_X.head()
```

Out[18]:

	url_length	hostname_length	path_length	FirstDir_length	total-	total@	total?	total%	total.
0	22	14	0	0	0	0	0	0	2
1	23	15	0	0	0	0	0	0	2
2	24	16	0	0	0	0	0	0	2
3	21	13	0	0	0	0	0	0	2
4	25	17	0	0	0	0	0	0	2

In [19]:

```
df_ml_train_Y = df_ml_train['result']
```

In [20]:

```
from imblearn.over_sampling import SMOTE
oversample = SMOTE()
```

In [21]:

```
x_sample_train, y_sample_train = SMOTE().fit_resample(df_ml_train_X, df_ml_train_Y.v
```


In [22]:

```
print(df_ml_train_X.shape)
print(df_ml_train_X.shape)
print(x_sample_train.shape)
print(y_sample_train.shape)
```

```
(450176, 20)
(450176, 20)
(691476, 20)
(691476,)
```

In [23]:

```
#finalize test dataset
```

In [24]:

```
df_test= pd.read_csv("URL_test.csv")
```

In [25]:

```
df_test.head()
```

Out[25]:

alhttps	totalwww	totaldigits	totalletters	totaldirs	totalnetlocdots	totalTLD	use_of_ip	short_url
0	0	18	49	4	1	1	1	1
0	0	23	43	4	1	1	1	1
0	0	22	44	4	1	1	1	1
0	0	18	46	4	1	1	1	1
0	0	18	48	4	1	1	1	1

In [26]:

```
df_test.drop(['Unnamed: 0', 'url', 'label'], axis=1, inplace =True)
```

In [27]:

```
df_test_X = df_test.loc[:, df_test.columns != 'result']
```

In [28]:

```
df_test_Y = df_test['result']
```

In [29]:

```
print(df_test_X.shape)
print(df_test_Y.shape)
```

```
(143901, 20)
(143901,)
```

KNN with Hyper-parameter tuning

In [30]:

```
from sklearn.neighbors import KNeighborsClassifier
```

In [31]:

```
knn_model_hyper = KNeighborsClassifier(n_neighbors=10, weights='distance', algorithm='
                                     leaf_size=30, p=1, metric='manhattan')
```

In [32]:

```
knn_model_hyper.fit(x_sample_train, y_sample_train)
```

Out[32]:

```

KNeighborsClassifier
KNeighborsClassifier(algorithm='ball_tree', metric='manhattan', n_nei
ghbors=10,
                    p=1, weights='distance')
```

In [33]:

```
knn_pred = knn_model_hyper.predict(df_test_X)
```

In [34]:

```
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
```

In [36]:

```
knn_accuracy = accuracy_score(df_test_Y, knn_pred)
```

In [37]:

```
print("Accuracy of KNN (Hyper-parameter) on Testing dataset : ", round(knn_accuracy, 3))
```

```
Accuracy of KNN (Hyper-parameter) on Testing dataset : 0.437
```

KNN : without hyper-parameter tuning

In [38]:

```
knn_model_2 = KNeighborsClassifier()
```

In [39]:

```
knn_model_2.fit(x_sample_train,y_sample_train)
```

Out[39]:

```
▼ KNeighborsClassifier  
KNeighborsClassifier()
```

In [40]:

```
knn_pred_2 = knn_model_2.predict(df_test_X)
```

In [41]:

```
knn_accuracy_2 = accuracy_score(df_test_Y, knn_pred_2)
```

In [42]:

```
print("Accuracy of KNN (without Hyper-parameters) on Testing dataset : ",round(knn_a
```

Accuracy of KNN (without Hyper-parameters) on Testing dataset : 0.445

Decision Tree without Hyper-parameter tuning

In [43]:

```
from sklearn.tree import DecisionTreeClassifier
```

In [44]:

```
dt_model = DecisionTreeClassifier()
```

In [45]:

```
dt_model.fit(x_sample_train, y_sample_train)
```

Out[45]:

```
▼ DecisionTreeClassifier  
DecisionTreeClassifier()
```

In [46]:

```
dt_model_pred = dt_model.predict(df_test_X)
```

In [47]:

```
dt_accuracy = accuracy_score(df_test_Y, dt_model_pred)
```

In [48]:

```
print('The accuracy of Decision Tree Classifier is : ', round(dt_accuracy,3))
```

The accuracy of Decision Tree Classifier is : 0.753

Decision Tree with Hyper-parameter tuning

In [49]:

```
dt_model_hyper = DecisionTreeClassifier(criterion='entropy')
```

In [50]:

```
dt_model_hyper.fit(x_sample_train, y_sample_train)
```

Out[50]:

```
DecisionTreeClassifier
DecisionTreeClassifier(criterion='entropy')
```

In [51]:

```
dt_model_pred_hyper = dt_model_hyper.predict(df_test_X)
```

In [52]:

```
dt_accuracy_hyper = accuracy_score(df_test_Y, dt_model_pred_hyper)
```

In [53]:

```
print('The accuracy of Decision Tree Classifier (with hyper-parameter tuning) is : '
```

The accuracy of Decision Tree Classifier (with hyper-parameter tuning)
is : 0.753

Random Forest without Hyper-parameters

In [54]:

```
from sklearn.ensemble import RandomForestClassifier

rf_model = RandomForestClassifier()

rf_model.fit(x_sample_train, y_sample_train)

rf_model_pred = rf_model.predict(df_test_X)

rf_accuracy = accuracy_score(df_test_Y, rf_model_pred)

print('The accuracy of Random Forest(without hyper-parameer) is : ', round(rf_accu
```

The accuracy of Random Forest(without hyper-parameer) is : 0.754

Random Forest with Hyper-parameters

In [55]:

```
from sklearn.ensemble import RandomForestClassifier

rf_model_hyper = RandomForestClassifier(n_estimators=150,max_features=None, criterion='entropy')
rf_model_hyper.fit(x_sample_train, y_sample_train)

rf_model_hyper_pred = rf_model_hyper.predict(df_test_X)

rf_hyper_accuracy = accuracy_score(df_test_Y, rf_model_hyper_pred)

print('The accuracy of Random Forest(with hyper-parameteer tuning) is : ', round(rf_hyper_accuracy, 3))

The accuracy of Random Forest(with hyper-parameteer tuning) is : 0.754
```

In [56]:

```
print(rf_hyper_accuracy)

0.7543172041889911
```

Extremely Randomized Tree Classifier

In [58]:

```
from sklearn.ensemble import ExtraTreesClassifier

xrf_model = ExtraTreesClassifier()
xrf_model.fit(x_sample_train, y_sample_train)
xrf_pred = xrf_model.predict(df_test_X)
xrf_accuracy = accuracy_score(df_test_Y, xrf_pred)
print("Accuracy of XRF is : ", round(xrf_accuracy, 3))

Accuracy of XRF is : 0.754
```

SAVING the model

In [60]:

```
import pickle
```

In [61]:

```
with open('saved_xrf', 'wb') as f:
    pickle.dump(xrf_model, f)
```

Trying models after scaling

In [63]:

```
from sklearn import preprocessing
```

In [64]:

```
scaling = preprocessing.StandardScaler()
```

In [65]:

```
x_sample_train_scaled = scaling.fit_transform(x_sample_train)
```

In [66]:

```
df_test_X_scaled = scaling.fit_transform(df_test_X)
```

In [68]:

```
x_sample_train_scaled
```

Out[68]:

```
array([[ -0.85279654,  -0.64418721,  -1.05736092, ...,   2.14813952,
         0.11976771,   0.26375863],
       [ -0.83153122,  -0.51045365,  -1.05736092, ...,   2.14813952,
         0.11976771,   0.26375863],
       [ -0.8102659 ,  -0.3767201 ,  -1.05736092, ...,  -0.26136056,
         0.11976771,   0.26375863],
       ...,
       [ -0.78900058,  -1.31285499,  -0.74631482, ...,  -0.26136056,
         0.11976771,  -4.00865115],
       [ -0.2998982 ,  -0.91165432,   0.07018117, ...,  -0.26136056,
         0.11976771,   0.26375863],
       [  0.82716379,   0.15821412,   1.70317316, ...,  -0.26136056,
         0.11976771,   0.26375863]])
```

Extra Tree after scaling

In [70]:

```
xrf_model_scaled = ExtraTreesClassifier(criterion='entropy')
xrf_model_scaled.fit(x_sample_train_scaled, y_sample_train)
xrf_pred_scaled = xrf_model_scaled.predict(df_test_X_scaled)
xrf_scaled_accuracy = accuracy_score(df_test_Y, xrf_pred_scaled)
print("Accuracy of XRF is : ",round(xrf_scaled_accuracy,3))
```

Accuracy of XRF is : 0.747

XGBoost on scaled data

In [71]:

```
from xgboost import XGBClassifier
```

In [92]:

```
boost = XGBClassifier(n_estimators=250, max_depth=5, learning_rate=1, objective='bir
```

In [93]:

```
boost.fit(x_sample_train_scaled, y_sample_train)
```

Out[93]:

```
XGBClassifier(base_score=0.5, booster='gbtree', callbacks=None,
              colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1,
              early_stopping_rounds=None, enable_categorical=False,
              eval_metric=None, feature_types=None, gamma=0, gpu_id=-1,
              grow_policy='depthwise', importance_type=None,
              interaction_constraints='', learning_rate=1, max_bin=256,
              max_cat_threshold=64, max_cat_to_onehot=4, max_delta_step=0,
```

In [94]:

```
boost_pred_scaled = boost.predict(df_test_X_scaled)
```

In [95]:

```
boost_pred_scaled_accuracy = accuracy_score(df_test_Y, boost_pred_scaled)
```

In [97]:

```
print("Accuracy of XGBoost is : ",round(boost_pred_scaled_accuracy,3))
```

Accuracy of XGBoost is : 0.765

DEEP LEARNING

In [99]:

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, Activation
from tensorflow.keras.callbacks import ModelCheckpoint
from tensorflow.keras.callbacks import ReduceLROnPlateau
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Conv1D, MaxPooling1D, Flatten, Dropout, Embedding
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.callbacks import ModelCheckpoint
```

2022-11-29 21:06:51.449844: I tensorflow/core/platform/cpu_feature_guard.cc:193] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: AVX2 FMA

To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.

In [101]:

```
x_sample_train_scaled.shape
```

Out[101]:

```
(691476, 20)
```

In [102]:

```
# model building
model = Sequential()
model.add(Dense(32, activation = 'relu', input_shape = (20, )))
model.add(Dense(16, activation='relu'))
model.add(Dense(8, activation='relu'))
model.add(Dense(1, activation='sigmoid'))

model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 32)	672
dense_5 (Dense)	(None, 16)	528
dense_6 (Dense)	(None, 8)	136
dense_7 (Dense)	(None, 1)	9

=====
 Total params: 1,345
 Trainable params: 1,345
 Non-trainable params: 0
 =====

In [103]:

```
# compiling model
opt = keras.optimizers.Adam(learning_rate=0.0001)
model.compile(optimizer= opt ,loss='binary_crossentropy',metrics=['acc'])
```

In [105]:

```
# custom callback to stop the training when certain metric value is reached

# stop training when validation loss reach 0.1
class myCallback(keras.callbacks.Callback):
    def on_epoch_end(self, epoch, logs={}):
        if(logs.get('val_loss')>0.1):
            print("\nReached 0.1 val_loss so cancelling training!")
            self.model.stop_training = True

callback = myCallback()
```


In [106]:

```
ze=260, callbacks=[callback],validation_data=(df_test_X_scaled,df_test_Y),verbose=1)
```

Epoch 1/3

2657/2660 [=====>.] - ETA: 0s - loss: 0.1141 -

acc: 0.9745

Reached 0.1 val_loss so cancelling training!

2660/2660 [=====] - 7s 2ms/step - loss: 0.114

0 - acc: 0.9745 - val_loss: 1.8937 - val_acc: 0.4041

In [108]:

```
import matplotlib.pyplot as plt
```

In [109]:

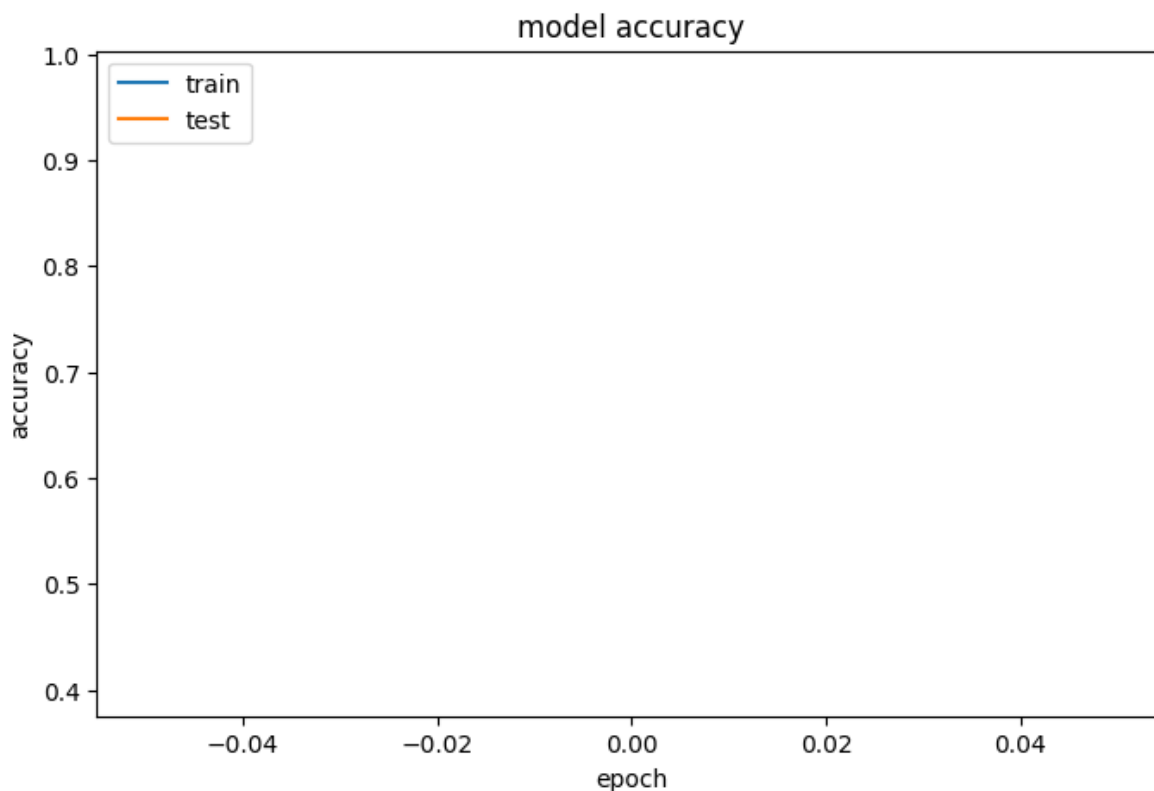
```
# DISPLAYING MODEL TRAINING HISTORY


# list all data in history
print(history.history.keys())

# summarize history for accuracy
plt.figure(figsize=(8,5))
plt.plot(history.history['acc'])
plt.plot(history.history['val_acc'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()

# summarize history for loss
plt.figure(figsize=(8,5))
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```

```
dict_keys(['loss', 'acc', 'val_loss', 'val_acc'])
```





In [110]:

```
# Evaluate the model on test dataset
loss, acc = model.evaluate(df_test_X_scaled, df_test_Y, verbose=1)
print('Test loss: {}'.format(loss))
print('Test Accuracy: {}'.format(acc))
```

```
4497/4497 [=====] - 3s 708us/step - loss: 1.8
937 - acc: 0.4041
Test loss: 1.8936975002288818
Test Accuracy: 0.40408337116241455
```

In [112]:

```
import numpy as np
```

In [113]:

```
# predicting on test data.
pred_test = model.predict(df_test_X_scaled)
for i in range (len(pred_test)):
    if (pred_test[i] < 0.5):
        pred_test[i] = 0
    else:
        pred_test[i] = 1
pred_test = pred_test.astype(int)
```

4497/4497 [=====] - 3s 616us/step

PREDICTED :

Mallicious
 Mallicious
 Mallicious
 Mallicious
 Mallicious
 Mallicious
 Mallicious
 Mallicious
 Mallicious
 Mallicious
 Mallicious
 Mallicious

ACTUAL :

Non Mallicious
 Non Mallicious
 Non Mallicious
 Non Mallicious
 Non Mallicious
 Non Mallicious
 Non Mallicious
 Non Mallicious
 Non Mallicious
 Non Mallicious
 Non Mallicious
 Non Mallicious

In []:

```
def view_result(array):
    array = np.array(array)
    for i in range(len(array)):
        if array[i] == 0:
            print("Non Mallicious")
        else:
            print("Mallicious")

print("PREDICTED : ")
view_result(pred_test[:11])
print("\n")
print("ACTUAL : ")
view_result(df_test_Y[:11])
```

NEW TENSOR FLOW

In [114]:

```
import tensorflow as tf
```

In [115]:

```
tf.random.set_seed(42)

model = tf.keras.Sequential([
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(256, activation='relu'),
    tf.keras.layers.Dense(256, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])

model.compile(
    loss=tf.keras.losses.binary_crossentropy,
    optimizer=tf.keras.optimizers.Adam(lr=0.03),
    metrics=[
        tf.keras.metrics.BinaryAccuracy(name='accuracy'),
        tf.keras.metrics.Precision(name='precision'),
        tf.keras.metrics.Recall(name='recall')
    ]
)

history = model.fit(x_sample_train_scaled, y_sample_train, epochs=100)

Epoch 95/100
21609/21609 [=====] - 37s 2ms/step - loss: 0.0109 - accuracy: 0.9977 - precision: 0.9976 - recall: 0.9978
Epoch 96/100
21609/21609 [=====] - 39s 2ms/step - loss: 0.0093 - accuracy: 0.9977 - precision: 0.9976 - recall: 0.9978
Epoch 97/100
21609/21609 [=====] - 38s 2ms/step - loss: 0.0107 - accuracy: 0.9977 - precision: 0.9976 - recall: 0.9978
Epoch 98/100
21609/21609 [=====] - 38s 2ms/step - loss: 0.0107 - accuracy: 0.9977 - precision: 0.9976 - recall: 0.9978
Epoch 99/100
21609/21609 [=====] - 39s 2ms/step - loss: 0.0106 - accuracy: 0.9977 - precision: 0.9977 - recall: 0.9978
Epoch 100/100
21609/21609 [=====] - 38s 2ms/step - loss: 0.0097 - accuracy: 0.9977 - precision: 0.9977 - recall: 0.9978
```


In [116]:

```
import matplotlib.pyplot as plt
from matplotlib import rcParams

rcParams['figure.figsize'] = (18, 8)
rcParams['axes.spines.top'] = False
rcParams['axes.spines.right'] = False
```

In [117]:

```
plt.plot(np.arange(1, 101), history.history['loss'], label='Loss')
plt.plot(np.arange(1, 101), history.history['accuracy'], label='Accuracy')
plt.plot(np.arange(1, 101), history.history['precision'], label='Precision')
plt.plot(np.arange(1, 101), history.history['recall'], label='Recall')
plt.title('Evaluation metrics', size=20)
plt.xlabel('Epoch', size=14)
plt.legend();
```



In [118]:

```
predictions = model.predict(df_test_X_scaled)
predictions
```

4497/4497 [=====] - 4s 828us/step

Out[118]:

```
array([[1.0000000e+00],
       [1.0000000e+00],
       [1.0000000e+00],
       ...,
       [7.4329654e-10],
       [7.5284329e-10],
       [1.0000000e+00]], dtype=float32)
```

In [119]:

```
prediction_classes = [1 if prob > 0.5 else 0 for prob in np.ravel(predictions)]
```

In [127]:

df_test_X_scaled

Out[127]:

```
array([[ -0.22924993, -1.41252652,  0.83056877, ..., -0.39404494,
         0.01419747,  0.20019433],
       [ -0.22924993, -1.41252652,  0.83056877, ..., -0.39404494,
         0.01419747,  0.20019433],
       [ -0.22924993, -1.41252652,  0.83056877, ..., -0.39404494,
         0.01419747,  0.20019433],
       ...,
       [ -0.31900902, -0.12666924, -0.64084925, ..., -0.39404494,
         0.01419747,  0.20019433],
       [ -0.65560562, -0.2874014 , -0.64084925, ..., -0.39404494,
         0.01419747,  0.20019433],
       [ -0.88000335,  3.7309026 , -0.86295009, ..., -0.39404494,
         0.01419747, -4.99514643]])
```

In [120]:

```
loss, accuracy, precision, recall = model.evaluate(df_test_X_scaled, df_test_Y)
loss, accuracy, precision, recall
```

```
4497/4497 [=====] - 5s 995us/step - loss: 74
8.7009 - accuracy: 0.6412 - precision: 0.7276 - recall: 0.8380
```

Out[120]:

```
(748.7008666992188, 0.6411769390106201, 0.7275629043579102, 0.83798825
74081421)
```

In [122]:

```
from sklearn.metrics import confusion_matrix

print(confusion_matrix(df_test_Y, prediction_classes))
```

```
[[ 1325 34053]
 [17582 90941]]
```

In [123]:

```
from sklearn.metrics import accuracy_score, precision_score, recall_score

print(f'Accuracy: {accuracy_score(df_test_Y, prediction_classes):.2f}')
print(f'Precision: {precision_score(df_test_Y, prediction_classes):.2f}')
print(f'Recall: {recall_score(df_test_Y, prediction_classes):.2f}')
```

```
Accuracy:  0.64
Precision: 0.73
Recall:    0.84
```

In []:

